

BONUS EXERCISE

– Experimenting with Deep Neural Networks –

Due: **Friday, November 16, 2018**, at 2:00 PM in Yaser’s office (150 Moore)

This is an optional exercise. If you decide to do it, and you want to get the bonus for it, you need to complete it and turn it in by the above due date. The exercise builds on topics in Lectures 10-12.

(prepared by *Aadyot Bhatnagar*)

0 Introduction

In this exercise, you will be exploring different neural net architectures. If you wanted to do this in the past, you would have been forced to implement your own neural net library. However, with the recent explosion of interest in deep learning,¹ many high quality open source implementations have proliferated, including TensorFlow, Torch, Caffe, CNTK, Theano, and Keras, to name just a few. Moreover, many of these libraries can run your deep learning code on the GPU instead of the CPU with no extra work needed on the part of the developer.

While implementing the forwards and backwards propagation algorithms yourself is a good exercise, that is not the goal of this assignment. Rather, you will be exploring the impacts of different architectural choices (different numbers of layers, hidden units, types of regularization, etc.) on the performance of a neural net. Accompanying this mini-project is a support code base written in Python that uses Keras to implement a neural network that recognizes handwritten digits (Figure 1). Most of your work will be focused on modifying the training script.

¹Deep learning is synonymous with neural networks, and the number of layers to be considered deep can be as little as 2 hidden layers ($L = 3$).

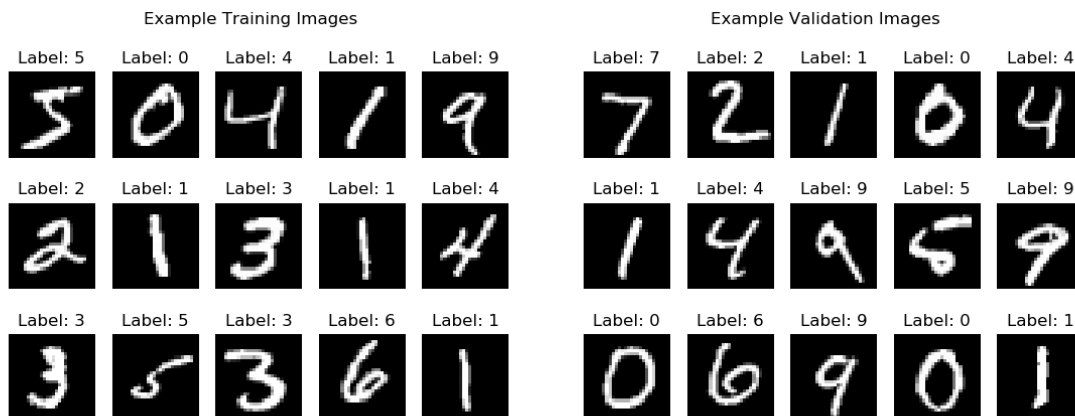


Figure 1: Example images from the MNIST dataset of handwritten digits.

We suggest that you use the `pip` package manager for Python 3 to run the code on your personal machine. This should come as a standard part of any Python 3 release. The support code has the following dependencies: `numpy`, `matplotlib`, `tensorflow`, `keras`, `simplejson`, `pillow`. To install a package `<pack>`, simply type `pip install <pack>` at the command line.

The support code is broken up into 3 scripts: `visualize.py`, `train.py`, and `evaluate.py` (in [this archive](#)). Use the `-h` command line option to see each of their usages. This is to reflect the typical workflow in machine learning and data science: you begin by exploring the type of data you have (images versus financial data, dimensionality, etc.), use that information to determine a model class to train on the data, and then evaluate the performance of that model. This is the way people typically modularize their code in industry. Why? If you decide that you need to make modifications to one part of the pipeline, you don't want to have to re-run everything else (which may be very time-consuming, as you will soon find out) to get new results.

Finally, please refer to the [Keras documentation](#) to figure out what specific calls in the support code are doing. This will be good practice for learning how to pick up new tools in the future.

1 CPU vs. GPU (Optional)

Complete this section only if you have access to a GPU machine, or if you want to set up a GPU instance on AWS. Note that with your `.edu` email address, you are eligible for free AWS credits! This would be a good time to practice using the cloud, as it's becoming more and more prevalent in industry. If you choose to use AWS, we suggest setting up an EC2 GPU instance with the Conda Deep Learning AMI, as described [here](#). You can then SSH into this machine and use it for anything you need.

If you choose to go this route, the only differences to the package installation instructions will be for installing TensorFlow and Keras. To do so, you would type

`conda install -c anaconda tensorflow-gpu` and `conda install keras` at the command line. However, with the Deep Learning AMI, it shouldn't be necessary to install TensorFlow, only Keras. You can also find a step-by-step guide [here](#).

We will not be supporting you to get GPU code running on your personal machine. However, if you are curious, you need to first install [CUDA 9.0](#) (NVIDIA's proprietary GPU programming language) and the associated library for deep learning, [cuDNN 7.3](#). This isn't too hard on Windows, but if you're working on Linux, this **may brick your machine!** Once this is done, you can just `pip install tensorflow-gpu`. The `tensorflow-gpu` package works as a Python front-end wrapper that invokes CUDA calls on the GPU behind the scenes. If you want to learn how to use CUDA and accelerate your programs using the GPU for other applications, consider taking [CS 179](#).

Once you've set up a GPU machine, run the training script for the dense² and convolutional³ neural networks given, both on a CPU on your personal machine and on the GPU machine provided. Note that you don't have to make any modifications to the source code to do this! How long does training take on each device?

2 Number of Parameters

Look at the neural net architecture specified by the `build_dense_net()` function in `train.py`. This is a deep neural net with 2 hidden layers, the first with 200 units, and the second with 100. We use weight-decay regularization⁴ on the learned weights matrix (more on this later), and follow each layer with a ReLU nonlinearity (also known as an activation; $ReLU(s) = \max\{0, s\}$ for any signal $s \in \mathbb{R}$) instead of $\tanh()$. The final dense layer is just to take a softmax⁵ and produce a probability distribution over $\{0, 1, \dots, 9\}$ (the possible labels we can assign an image).

Leave the architecture of this model fixed for now. What happens to the learning curve when you vary the number of hidden units? Specify the number of hidden units you use in each layer, and report learning curves for each different architecture you try. What trends do you notice? What is the smallest number of model parameters (not hidden units!) for which you can achieve over 95% validation accuracy? 97% validation accuracy?

3 Regularization

In this document, we will use 'loss' as synonymous with error. Hopefully, you saw that the huge neural net implemented in the support code achieves much lower training

²*Dense* means fully connected from layer to layer, as the neural-network model given in class.

³*Convolutional* is a special neural-network model explained later.

⁴Also known as ℓ^2 regularization since it uses the 2-norm of \mathbf{w} .

⁵*Softmax* is a generalization of the logistic function to multiple variables; $\theta_i(s_1, \dots, s_K) = e^{s_i} / \sum_{k=1}^K e^{s_k}$ which transforms signals into probabilities.

loss than validation loss (validation loss is the error on a hold-out dataset that we use as an estimate for out-of-sample performance). While the MNIST dataset is a toy problem that many models have an easy time performing well on, this disparity is the kiss of death for machine learning systems used on harder real-world tasks (like predicting financial markets and planning autonomous navigation, to name two). Specifically, this difference is a symptom of a complex model overfitting to the training data. The whole reason we have a hold-out dataset used exclusively for validation is to get a feel for how well our model generalizes out of sample.

One way to address this issue is through *regularization* (Lecture 12). Think of it as follows: instead of having our model minimize the training loss $L(\mathbf{w} \mid x, y)$ (which is the in-sample error of the hypothesis defined by \mathbf{w}) in a vacuum, we also penalize it for being too complex. The most straightforward way to do so is to minimize an *augmented* loss function $L(\mathbf{w} \mid x, y) + \lambda \|\mathbf{w}\|^2$ instead of the loss $L(\mathbf{w} \mid x, y)$ on its own.

To gain some intuition for this augmented loss, consider the simpler example of polynomial regression, where our model class is $f(x \mid \mathbf{w}) = w_0 + w_1x + \dots + w_kx^k$. If we have large coefficients w_i , then we can have a highly oscillatory polynomial, which is unlikely to represent the underlying data. By forcing $\|\mathbf{w}\|^2$ to be small, we reduce the model complexity. See Figure 2 for a concrete example of minimizing $L(\mathbf{w} \mid x, y) + \lambda \|\mathbf{w}\|^2$ (where $L(\mathbf{w} \mid x, y)$ is mean squared error) for different regularization strengths λ .




Figure 2: Fitting a 9th order polynomial regression to data generated noisily from a cubic function. Small λ results in overfitting to the noise, intermediate λ results in good generalization, and large λ reduces model complexity too much to be useful.

So in this problem, explore the usefulness of weight-decay regularization (minimizing $L(\mathbf{w} \mid x, y) + \lambda \|\mathbf{w}\|^2$ instead of just $L(\mathbf{w} \mid x, y)$) on the large neural net implementation given in the support code. You can do so by specifying the `-r` command line option in `train.py`. Report learning curves for different choices of λ . Can you use regularization to get a similar result to one of the smaller neural nets you implemented in the previous part?

4 Activations

In class, we have primarily discussed using the tanh and sigmoid activations. In practice nowadays, people often use the ReLU activation, which is defined as $ReLU(s) = \max\{0, s\}$. We use this activation function in the support code, and in practice, people have observed that it results in much faster convergence. Verify this for yourself by replacing the ReLU activation with a tanh activation. Why do you think this happens? (Hint: how does $\tanh(s)$ behave as $|s| \rightarrow \infty$? Compare this to $ReLU(s)$.)

5 Different Architectures

Now, let's say that I give you a fixed budget of 200 hidden units. What is the best validation accuracy you can achieve? Feel free to vary the number of layers, the kind of regularization (e.g., ℓ^1 which is based on the 1-norm of \mathbf{w} as opposed to ℓ^2 which is the usual weight decay based on the 2-norm) and its strength (λ), the activation, and the optimizer you use.

6 Convolutional Neural Nets

Finally, we will discuss a different sort of regularization, known as an *inductive bias*. Inherently, a densely connected neural net is not a model class optimized for image data. It is unrealistic to model a connection between every pixel in an image. Dependencies in image data are much more local, i.e. adjacent pixels are much more likely to be correlated with each other than those found at opposite ends of an image. Moreover, images have spatial invariance: a face is a face, whether we find it at the top left, center, or bottom right of an image.

Without going too much into the specific details (you can learn more about this in [CS 155](#)), a 2D convolution satisfies all these properties much better than a matrix multiplication. As a result, by learning a few smaller convolutional filters instead of a series of huge matrices, a convolutional neural net performs much better on this image classification task with a fraction of the parameters. Verify this for yourself. Run `train.py` and `evaluate.py` using the convolutional neural net and no regularization. Compare the number of parameters between any of the convolutional and dense neural nets, and report learning curves.

7 Feedback

What did you think of this mini-project? What changes would you suggest we make to improve this for future iterations of the class?